

# **70042 Group Project Report: Memory Allocation for Kernel-Bypass Frameworks**

Ethan Sim (ers19)

Alvin Jonathan (aj22)

Brendan Cheu (tcc22)

Christopher Adnel (ca222)

Si Wei Liang (sl222)

**Supervisor:** Marios Kogias

**Date:** 9th March 2022

## **1. Background**

Caching systems store responses to frequently-made requests from users within main memory [1]. This reduces expensive resource usage (e.g., backend databases), which lowers latency [2]. Many popular web applications, such as Facebook and Wikipedia, use Memcached, a key-value store-based caching system termed “the gold standard of web caching” [3]. Despite its popularity, Memcached suffers from well-documented performance limitations. Notably, when Memcached is run on vanilla operating systems, packets are processed by these operating systems’ network stacks, and significant delays are incurred due to system calls [4]. As the number of system calls made increases with the number of packets received [5], aggregating requests into fewer packets of greater size will reduce these delays, improving performance.

However, if request aggregation involves the kernel when processing packets, this overhead is not absorbed, only transferred. Kernel-bypassing frameworks, such as the Data Plane Development Kit (DPDK) [6], overcome this by manipulating packets entirely within userspace. Yet, modifying existing applications to use these frameworks may be costly, as the application must now implement and attain high performance on a dedicated network stack [2]. We therefore implement a software middlebox on top of DPDK, which absorbs the packet-processing overhead without requiring the modification of deployed applications.

## **2. Project Architecture**

In the context of our project, *clients* are users seeking to retrieve values from Memcached, a key-value store. These clients issue `get` requests using User Datagram Protocol (UDP) in the format “**GET <key 1> <key 2>...**”. Requests contain at least one key, but, per the Memcached protocol, cannot exceed the length of a single UDP packet [7]. Upon receiving a `get` request, Memcached returns the corresponding value for each key in an individual UDP reply with the format “**VALUE <request id> <value length> <unique id> <actual value>**”, which may comprise multiple packets. Our *middlebox* sits between clients and Memcached: it will aggregate client requests before passing these aggregated requests on to Memcached. Upon receiving a reply from Memcached, it will steer the returned values from Memcached to the respective clients.

Firstly, the Middlebox receives the packet at the Network Interface Card (NIC). The Middlebox will bypass the kernel space to begin filtering and processing the packets at the user space. This filtering is conducted by looking at the headers of the packet. Only packets that use Internet Protocol (IP), communicated via UDP and containing expected source and destination IP addresses and ports are retained.

Secondly, the Memcached UDP protocol frame header (not to be confused with the UDP header of the transport layer) at the start of the payload segment of the packet will be processed. The frame header is 8 bytes long, and only the first 2 bytes, which represents the client’s request id, will be used by the Middlebox.

Thirdly, the Middlebox will process the payload to extract the keys in the client’s request and merge each key with a new request ID. The new request ID is generated by the Middlebox to identify each request made by the Middlebox to Memcached. The result of the merging will be used as separate keys for the

hashmap. The hashmap stores this newly generated key and the client info to keep a record of the clients that expects a reply for each key in each Memcached request. Concurrently, a unique key list, that filters out requests for the same key, is being populated.

Finally, when the unique key list fully fills a packet, the keys will be sent as a single request to Memcached. Middlebox implements an algorithm to handle low loads when we do not have enough keys to fill a packet. The Middlebox will populate headers of the outgoing network packet as the packet exits.

When the Memcached server returns with the values for each key separately, the hashmap will then be used to look for the relevant clients to forward each reply to. Figure 1 summarises our approach.

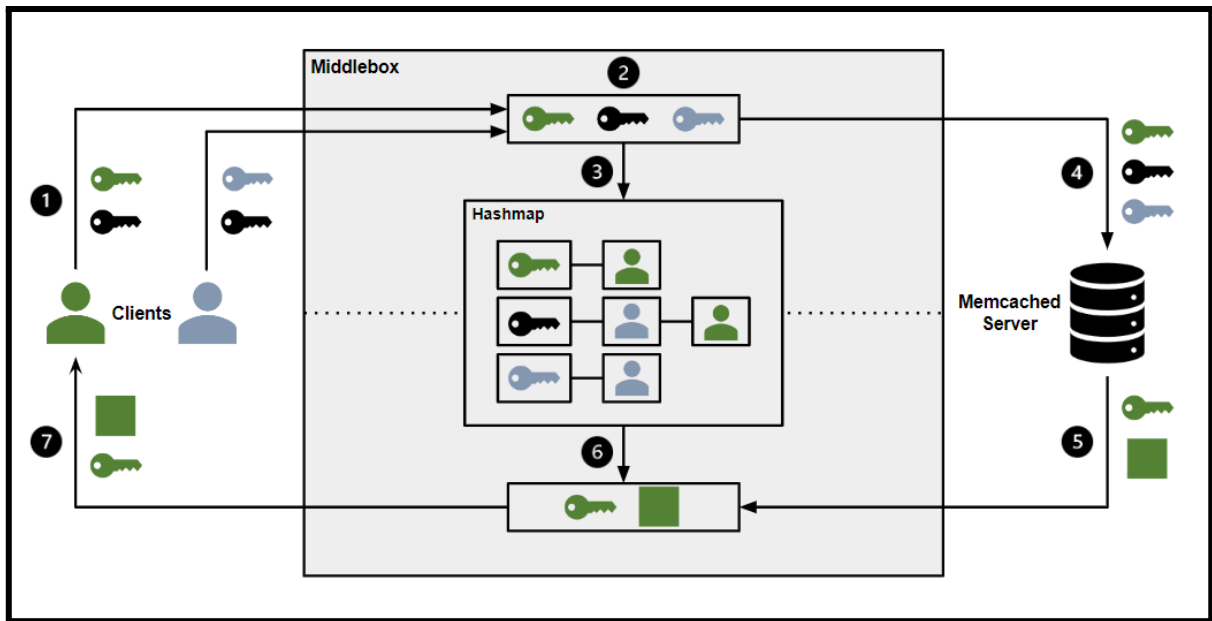


Figure 1: Summary of middlebox operation. Two clients each send a two-key get request to the middlebox. The black key appears in both clients' requests (1). Upon reaching the middlebox, the requested keys are extracted and combined into a single packet; duplicates are removed (2). Client information, such as the IP address, is concurrently extracted and associated with the requested keys using a hashmap (3). The combined packet is then sent to Memcached (4). Upon receiving a reply to a requested key containing the key and its associated value (5), the key is used to query the hashmap, identifying the requesting client (6), and the reply is forwarded accordingly (7).

### 3. Implementation

The development of the Middlebox took reference from the team supervisor's previous work to achieve the following functionalities [8]. Firstly the Middlebox, being a latency-critical component, achieves kernel-bypassing by using DPDK libraries as seen in figure 2.

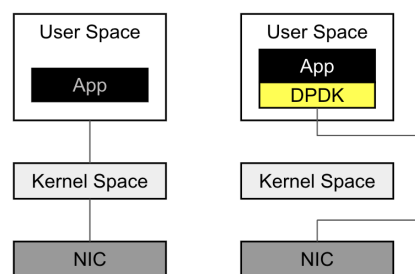


Figure 2: How applications communicate with NIC without (left) and with (right) DPDK

DPDK libraries allow the Middlebox to communicate with the NIC and obtain packets from the NIC's queue without using the kernel space, reducing overhead incurred due to system calls and context

switches. Additionally, using DPDK avoids system calls when memory allocation is performed during runtime [6]. DPDK manages its own pool of memory which is allocated at the start of the program. Therefore, memory allocation can be performed using the memory pool without the need to use functions, such as malloc(), that involve system calls. Memory is also reused whenever possible.

At low loads, there will not be enough requests from clients to fully fill a packet. Thus, an algorithm, seen in figure 3, is implemented to avoid a scenario where the client is waiting for too long or the Memcached is left idling while client requests are being processed. This is achieved by ensuring that the outstanding number of packets at the server is not too low and the timing each request stays in the Middlebox is not too long using a configurable timeout.

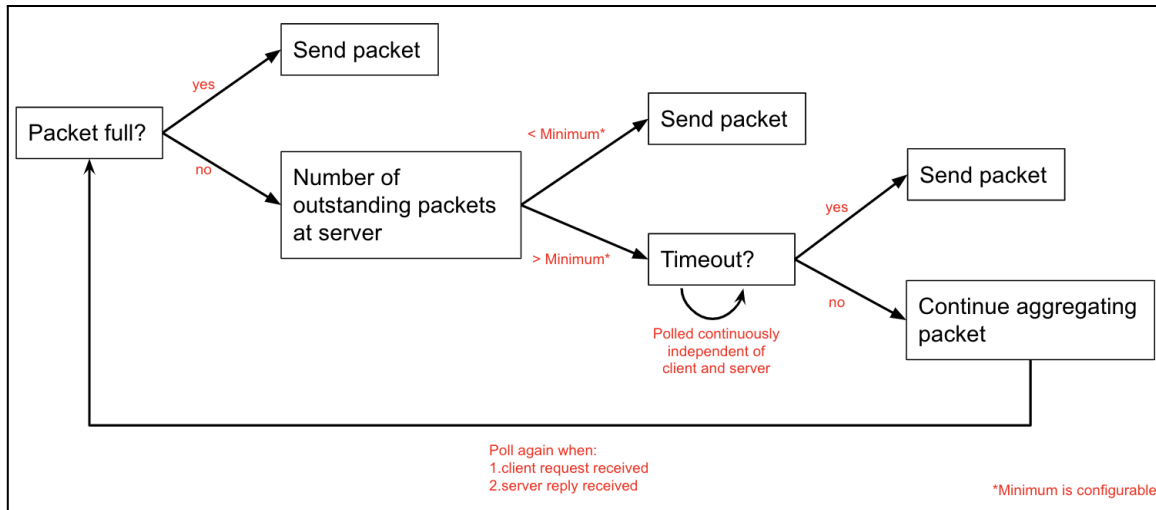


Figure 3: Algorithm implemented in Middlebox

#### 4. Evaluation

This section describes the tests performed to ensure that the expected functionalities of the Middlebox are met. Screenshots of Wireshark’s “Packet Detail Panes” and “Packet List Pane” show each packet’s content, and source and destination I.P respectively, in the same sequence, for all tests. The tests involves client 1 (10.1.0.1) and client 2 (10.1.0.4) communicating with Memcached (10.1.0.3) through the Middlebox (10.1.0.2).

The first test conducted ensured that the middlebox was correctly merging packets from multiple clients and follows the Algorithm in section 3. This can be observed in the wireshark screenshots in figure 4.

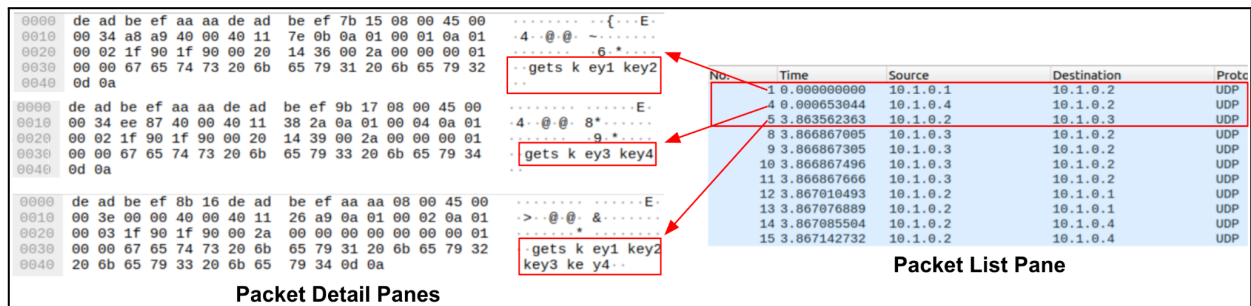


Figure 4: Client 1 sends requests for key1 and key 2 and client 2 sends requests for key 3 and key4. Middlebox successfully merges the packets and forwards one request (“gets key1 key2 key3 key4”) to server. Also, Middlebox follows the algorithm and sends to Memcached even though it is not fully filled as we are below the pre-set number of outstanding requests at Memcached. Middlebox is expected to produce the same result as above in a timeout as well.

The return path of replies from Memcached in the first test allowed us to ensure that the Middlebox is able to route the replies from Memcached to the right clients using the Hashmap. This can be observed in figure 5, which shows the values returned by Memcached to the multikey get request made in figure 4.

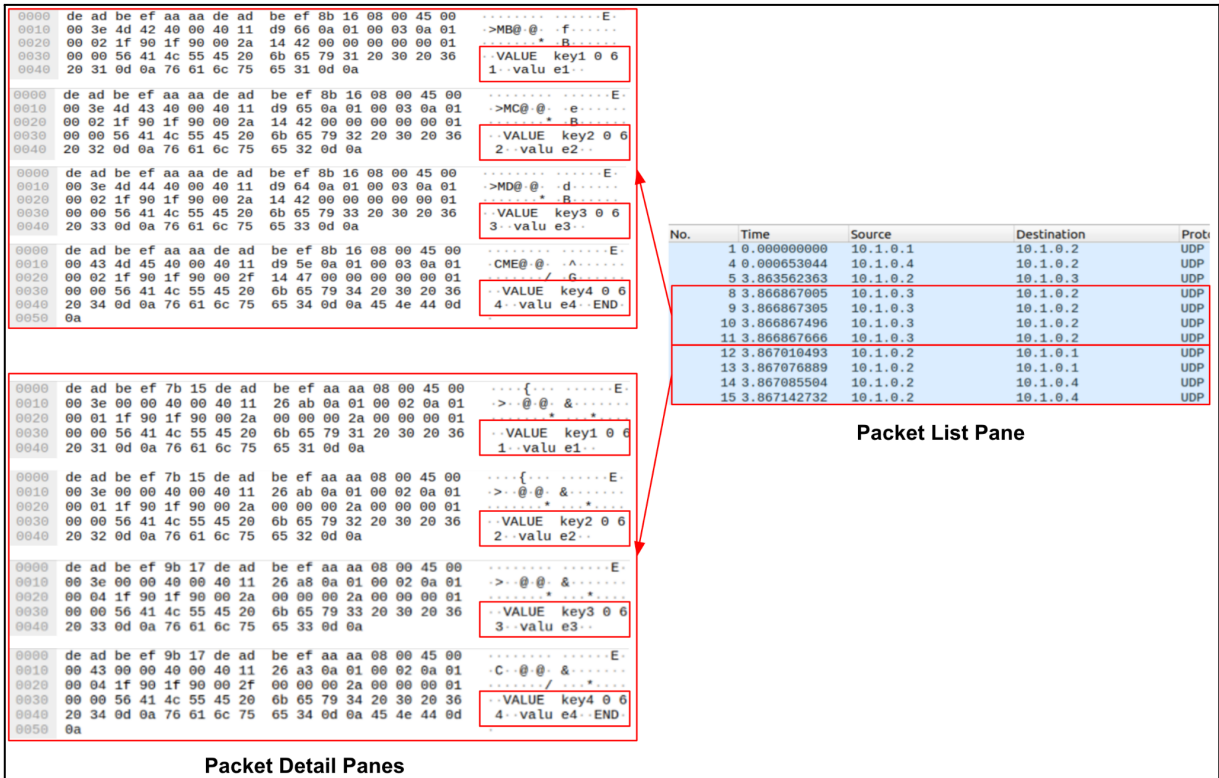


Figure 5: The top left image shows the values returned by Memcached for all 4 keys (in figure 4) separately. From the packet list pane, it can be observed that each reply in the bottom left image is being routed correctly to client 1 and 2.

The second test ensures that when a request containing a single key has a value that requires a multipacket reply from Memcached to complete, the Middlebox is able to route the replies to the correct client, as seen in figure 6.

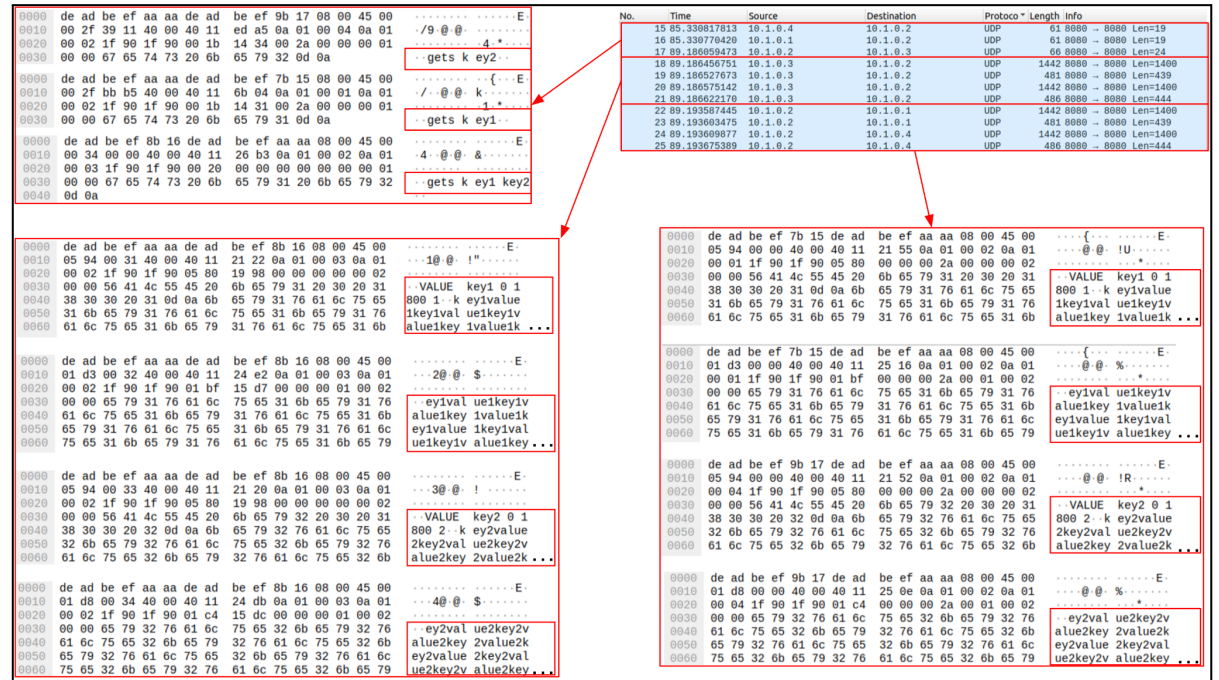


Figure 6: Using the information in the Packet List Pane with each image: **Top left image:** client 1 and 2 made a request for key1 and key2 respectively, Middlebox sends the merged request to Memcached. **Bottom left image:** Memcached replies with a value that exceeds 1 packet length for each key (key1 with key1value1key1... and key2

with key2value2key2...). **Bottom right image:** key1 is forwarded to client 1 (10.1.0.1) and key2 is forwarded to client 2 (10.1.0.4), indicating forwarding done correctly.

\* Details to set up, run and test the Middlebox with the Clients and Memcached can be found at:  
<https://github.com/ajonw/group-project>.

## 5. Limitations

The Middlebox does not guarantee that it is able to correctly pass to clients multipacket replies if in the same aggregated packet, 2 or more clients expect multipacket replies. When Memcached makes a multipacket reply, only the first packet contains information that the Middlebox uses to map back to the respective clients. Subsequent packets do not contain such information, thus the Middlebox is unable to identify which clients it should be forwarded to.

## 6. Future Work

First, by leveraging the existing Middlebox, future work could look into benchmarking the performance and quantifying the performance benefit of having the middlebox aggregate requests, by comparing it with the case where clients send requests directly to the server in the same environment. Lastly, future research may look into the possibility of enabling requests to be forwarded to multiple servers instead of just one. This will aid load balancing by spreading packets among multiple servers.

## 7. References

- [1] Mertz J, Nunes I. A qualitative study of application-level caching. IEEE Transactions on Software Engineering. 2016 Dec 1;43(9):798-816.
- [2] Ghigoff Y, Sopena J, Lazri K, Blin A, Muller G. BMC: Accelerating Memcached using Safe In-kernel Caching and Pre-stack Processing. In NSDI 2021 Apr 12 (Vol. 3, p. 4).
- [3] Wiger N, Imalsina R. Performance at Scale with Amazon ElastiCache San Diego, California: Amazon Web Services; 2019 p. 3–4.
- [4] Chidambaram V, Ramamurthi D. Performance analysis of memcached. unpublished Manuscript.[Online]. Available: <http://citeseerx.ist.psu.edu/viewdoc/download>. 2010.
- [5] Hruby T, Crivat T, Bos H, Tanenbaum AS. On Sockets and System Calls: Minimizing Context Switches for the Socket {API}. In 2014 Conference on Timely Results in Operating Systems ({TRIOS} 14) 2014.
- [6] Memory in DPDK, part 1: General concepts [Internet]. DPDK. LF Projects; 2019 [cited 2023 Mar 9]. Available from: <https://www.dpdk.org/blog/2019/08/21/memory-in-dpdk-part-1-general-concepts/>
- [7] Mate B. protocol.txt [Internet]. 2022. Available from: <https://github.com/memcached/memcached/blob/master/doc/protocol.txt>
- [8] Marios K. netstack [Internet]. 2020. Available from: <https://github.com/epfl-dcsl/r2p2/tree/master/netstack>